

Prathyush Sajith

psaji@uic.edu | prathyushsajith17@gmail.com | LinkedIn | Portfolio
Chicago, Illinois - 60612

EDUCATION

- **University of Illinois at Chicago** Jul 2026 (expected)
MS Thesis in Computer Engineering
 - Advisor: Prof. [Ahmet Enis Cetin](#)
 - Selected Courses: Convex Optimization, Neural Networks, Image Analysis and CV, Pattern Recognition.
- **SRM Institute of Science and Technology, Ramapuram** May 2024
B.Tech in Computer Science Engineering
 - Grade: 9.20/10.00 (First class with Distinction)

EXPERIENCE

- Graduate Research Assistant – Machine Learning & Data Analyst** Jan 2026 – Present
[Tools used: Python - Pandas, PyTorch, Scikit-learn, OpenCV, quPATH, Image] // HIPAA certified]
- Developed a bioimaging data pipeline using Pandas and Matplotlib to parse and analyze QuPath-exported CD31/Fibrinogen co-localization metrics across multiple ROIs to quantify regional blood-brain barrier permeability.
 - Inferred on a Mask R-CNN instance segmentation model via Amazon SageMaker JumpStart to automate detection and morphological analysis of cellular structures in high-resolution but noisy TIRFM imagery.
 - Used Pretrained models to reconstruct high-resolution 3D mappings of mouse brain structures from TIRFM images.

MS THESIS

- Improving Efficiency of Stereo Depth Estimation using Walsh-Hadamard Transforms** May 2025 – Present
[Skills used: PyTorch, TorchVision, OpenCV, CUDA, cuDNN, PIL, Matplotlib, Epipolar Geometry, Stereo Rectification]
- Designed and extended a **Stereo Transformer (STTR)** based stereo depth estimation pipeline on a custom **CARLA** synthetic dataset, leveraging attention-driven correspondence modeling with 3 times lesser parameters.
 - Improved computational efficiency by replacing standard Conv2D layers with a Walsh-Hadamard Transform-based convolution (WHTConv2D) **by 18.33%**, reducing complexity while preserving representational capacity.
 - Implemented a **Log-Disparity Loss formulation** to emphasize supervision on far-range (>100 m) depth regions, improving stability and accuracy for small-disparity predictions.

RECENT PROJECTS

- (1) UIC AI-Agent Hackathon - 2026 (git)** May 2026
[Tools used: Python, SQL, Cloudflare Workers, Cloudflare D1 (SQLite), Anthropic Claude API, Wrangler]
Built a full-stack resource kit for a UIC College of Business hackathon, enabling teams to develop AI agents that address real-world healthcare challenges using synthetic patient data (117 patients, \$27.9M in healthcare costs). Designed 5 tools with SQL queries to retrieve and perform tasks (ED risk detection, cost explanation, care barrier analysis). Implemented progressive workshop examples covering chat, tool use, and full agentic loops in Python.
- (2) History-Aware Local RAG System with LangChain & Ollama (git)** Apr 2026
[Tools used: Python, LangChain (LCEL), Ollama (Llama 3.2), HuggingFace Embeddings (all-MiniLM-L6-v2), FAISS]
Built a fully local Retrieval-Augmented Generation system to query about 3.4M characters of literary text with no external API dependencies. Chunked documents into 4,700+ segments, embedded them using HuggingFace's all-MiniLM-L6-v2, and indexed with FAISS for semantic search. Implemented MMR-based retrieval and a conversational interface with history-aware question rewriting using LangChain LCEL and a locally hosted Llama 3.2 via Ollama.
- (3) Custom built GPT-Based Large Language Model (git)** Jan – Feb 2025
[Tools used: Pytorch, Hugging Face Transformers, NumPy, Pandas]
Built a 124M parameter GPT-based LLM using Pytorch. Designed an efficient tokenizer, created embeddings with 768-1280 dimensions, and implemented multi-head self-attention and rotary embeddings. Pretrained the model on up to 1024 tokens per context window and achieved significant loss reduction. Fine-tuned to generate coherent text outputs.

SKILLS

- **Languages:** Python, SQL, MATLAB, (familiar with C and C++)
- **Deep Learning:** PyTorch, TorchVision, TensorFlow, Keras, JAX, scikit-learn, Hugging Face
- **3D & Computer Vision:** Camera calibration, instance segmentation, OpenCV, PIL
- **Model Optimization & Efficiency:** Spectral convolutions (Walsh-Hadamard / DCT), FLOPs / MACs profiling, latency benchmarking, torch.compile (Inductor), CUDA, cuDNN, TensorRT, ONNX Runtime
- **DL Architectures:** Transformers & attention (STTR), CNNs (PSMNet), U-Net, GANs, rotary embeddings
- **Generative AI / LLMs:** RAG, LangChain, FAISS, Ollama, Anthropic Claude API
- **MLOps & Cloud:** AWS (SageMaker, EC2, S3), Docker, Kubernetes, Weights & Biases, MLflow, Git, CI/CD
- **Analysis:** NumPy, Pandas, Matplotlib, Seaborn, Tableau, Power BI, EDA, data cleaning